

---

# Faster Neural Machine Translation Inference

**Hieu Hoang**

hieu@hoang.co.uk

**Tomasz Dwojak**

???

Adam Mickiewicz University

**Kenneth Heafield**

???

University of Edinburgh, Scotland

---

## Abstract

Deep learning models are widely deployed for machine translation. In comparison on many other deep learning models, there are two characteristics of machine translation which have not been adequately addressed by most software implementations, leading to slow inference speed. Firstly, as opposed to standard binary class models, machine translation models are used to discriminate between a large number of classes corresponding to the output vocabulary. Secondly, rather than a single label, the output from machine translation models is a sequence of class labels that makes up the words in the target sentence. The sentence lengths are unpredictable, leading to inefficiencies during batch processing. We provide solutions for these issues which together, increase batched inference speed by up to ??? on modern GPUs without affecting model quality. For applications where the use of maxi-batching introduces unacceptable delays in response time, our work speeds up inference speed by up to ???. Our work is applicable to other language generation tasks and beyond.

## 1 Introduction

We will look at two areas that are critical to fast NMT inference where the models in NMT differ significantly from those in other applications. These areas have been overlooked by the general deep-learning community, we aim to improve their efficiency for NMT-specific tasks.

Firstly, the number of classes in many deep-learning applications is small. However, the number of classes in NMT models is typically in the tens or hundreds of thousands, corresponding to the vocabulary size of the output language. For example, Sennrich et al. (2016) experimented with target vocabulary size of 60,000 and 90,000 subword units. This makes the output layer of NMT models very computationally expensive. ??? shows the breakdown of amount of time during translation in our NMT system; over 40% of the time is involved in calculating the activation, softmax and the beam search. We will look at optimizations which explicitly target the output layer.

Secondly, mini-batching is often used to increase the speed of deep-learning applications, including NMT systems. (??? graph of batch v. speed) shows using batching can increase speed by 17 times. However, mini-batching does not take into account the variable length of NMT inputs and outputs, creating computational issues and efficiency challenges. The computational issues are often solved by masking unwanted calculations. Translation speed can often be improved by employing max-batching whereby the input set is pre-sorted by sentence length before mini-batching, creating mini-batches with similar length input sentences. However, the target sentence lengths will still differ even for similar length inputs but the standard

mini-batching algorithm must continue to process the batch until all target sentences have been completed. (??? batchsize size v. iterations) shows the number of sentences still being processed for each decoding iteration, for a specific batch. The number of sentences still to be processed decreases, reducing the effectiveness of mini-batching. We will propose an alternative to the mini-batching algorithm.

We based our model on the sequence-to-sequence model of Cho et al. (2014) for machine translation models, but unlike Devlin (2017), we avoid solutions for the specific model. Therefore, our solution should be applicable to other models, architectures and task which have the similar characteristics. We envisage that our solutions would be of value to models used in text summarization, chatbot or image captioning. We also choose to focus on the use of GPU, rather than CPU as pursued in Devlin (2017).

## 2 Prior Work

Deep learning models have been successfully used on many machine learning task in recent years in areas as diverse as computer vision and natural language processing. This success have been followed by the rush to put these model into production. However, the computational resources required in order to run these models have ben chanllenging. One possible solution involves creating faster models that can approximate the original model (Kim and Rush, 2016). Another solution is tackling specific computationally intensive functions by approximating them (??? softmax).

Most current neural MT models follow an encoder-decoder architecture. The encoder calculates the word and sentence embeddings while the decoder generates the output sentence. The architecture within the encoder is still a subject of much research. Kalchbrenner and Blunsom (2013) used a convolution to encode the input and a recurrent neural network (RNN) for the decoder. RNNs was used for both encoding and decoding in (??? who used RNN 1st). (???) significantly improved translation quality by using LSTM was used in each RNN node. Cho et al. (2014) used GRU that has the required properties of LSTM but is computationally cheaper. Bahdanau et al. (2014) added attention model which improved translation at a cost of slower speed.

There has been recent attempts to move away from the sequence-to-sequence models of RNN encoder-decoder. Some justify their architecture on faster inference as well as better translation results. (??? Attention is all you need) and (??? coarse-to-fine) has been proposed as a fast alternatives to RNN-based models as it is possible to process more of the units in parallel.

It has been noticed that half precision arithmetic can be used for deep learning model without significant loss of model quality (Micikevicius et al., 2017). Other solutions include specialized hardware, the most popular being graphical processing units (GPU) but other hardware such as custom processors TPU (Jouppi et al., 2017), FPGA (Lacey et al., 2016) have been used.

Many of these optimization are general purpose improvements for deep learning models, regardless of the task it is being used in. Devlin (2017) is a noticeable machine translation-specific optimization which describe the speed improvements that can be obtained by techniques such as the use of half-precision, model changes and pre-computation.

## 3 Proposal

### 3.1 Softmax and Beam Search Fusion

The output layer of most deep learning models consist of the following steps

1. activation - multiplication of the input with the weight matrix for that layer

2. addition of a bias term to the resulting scores
3. activation function - in the output layers, this is commonly softmax or a variant, eg. log-softmax.
4. a search for the argmax output class, and probability is necessary.

In models with a small number of classes such as binary classification, the computational effort required is trivial and fast. However, this is not the case for large number of classes such as those found in NMT models.

We shall leave the matrix multiplication for the next proposal and future work, and concentrate on the last three steps, the outline of which are shown in Figures 1 to 3.

**Require:** activation vector  $p$ , bias vector  $b$   
**for all**  $p_i$  **in**  $p$  **do**  
     $p_i \leftarrow p_i + b_i$   
**end for**

Figure 1: Add Bias Term

**Require:** activation vector  $p$   
{calculate max for softmax stability}  
 $max \leftarrow -\infty$   
**for all**  $p_i$  **in**  $p$  **do**  
    **if**  $p_i > max$  **then**  
         $max \leftarrow p_i$   
    **end if**  
**end for**  
{calculate denominator}  
 $sum \leftarrow 0$   
**for all**  $p_i$  **in**  $p$  **do**  
     $sum \leftarrow sum + e^{p_i - max}$   
**end for**  
{calculate softmax}  
**for all**  $p_i$  **in**  $p$  **do**  
     $p_i \leftarrow \frac{e^{p_i - max}}{sum}$   
**end for**  
**return**  $p$

Figure 2: Calculate softmax

As can be seen, the operations iterate over the matrix  $p$  five times - once to add the bias, three times to calculate the softmax, and once to search for the best class. We shall fuse the three functions and combine the matrix access into two iterations, reducing memory bandwidth. Kernel fusion (Guevara et al., 2009) is a well practised optimization technique.

Secondly, we make use of the fact that softmax and exp are monotonic functions therefore we can avoid the search for the argmax in Figure 3 as it will be the same as that in Figure 2.

Thirdly, we are only interested in the probability of the best class. Since we have precalculated the best class in Figure 2, we shall only calculate softmax for this one class, avoiding the computational and memory bandwidth of computing the probability for all classes. The outline of our function is shown in Figure 4.

**Require:** softmax vector  $p$

```

 $max \leftarrow -\infty$ 
for all  $p_i$  in  $p$  do
  if  $p_i > max$  then
     $max \leftarrow p_i$ 
     $argmax \leftarrow i$ 
  end if
end for
return  $max, argmax$ 

```

Figure 3: Find best

**Require:** activation vector  $p$ , bias vector  $b$

```

{add bias, calculate  $max$  &  $argmax$ }
 $max \leftarrow -\infty$ 
for all  $p_i$  in  $p$  do
  if  $p_i + b_i > max$  then
     $max \leftarrow p_i + b_i$ 
     $argmax \leftarrow i$ 
  end if
end for
{calculate denominator}
 $sum \leftarrow 0$ 
for all  $p_i$  in  $p$  do
  if  $p_i > max$  then
     $sum \leftarrow sum + e^{p_i - max}$ 
  end if
end for
{calculate softmax( $p_i$ )}
return  $\frac{1}{sum}, argmax$ 

```

Figure 4: Combined softmax and argmax

In fact, we are usually only interested in the argmax class during inference, not the probability. Therefore, we can skip the second iteration over  $p$  and avoid computing the softmax altogether.

It has been shown (Koehn and Knowles, 2017) that using beam search to use the n-best number of classes, rather than just the best class, improves translation quality. This is a simple extension to the algorithm of Figure 4. Unlike the 1-best case, however, the softmax calculation cannot be skipped as the denominator differs for each input.

### 3.2 Top-up Batching

The standard mini-batching algorithm is outlined in Figure 5.

This algorithm encode the sentences for a batch, followed by decoding the batch. The decoding stop once all sentences in the batch are completed. This is a potential inefficiency as the number of remaining sentences may not be optimal.

We will focus on decoding as this is the more compute-intensive step, and issues with differing sentence sizes in encoding can partly be ameliorated by maxi-batching.

Our proposed top-up batching algorithm encode and decode asynchronously. The encod-

```

while more input do
  Create batch
  Encode
  while batch is not empty do
    Decode batch
    for each sentence in batch do
      if translation is complete then
        Remove sentence from batch
      end if
    end for
  end while
end while

```

Figure 5: Mini-batching

ing step, Figure 6, is similar to the main loop of the standard algorithm but the results are added to a queue to be consumed by the decoding step later.

```

while more input do
  Create encoding batch
  Encode
  Add to queue
end while

```

Figure 6: Encoding for top-up batching

Rather than decoding the same batch until all sentences in the batch are completed, the decoding step processes the same batch continuously. New sentences are added to the batch as old sentences complete, Figure 7.

```

create decoding batch  $b$  from queue
while  $b$  is not empty do
  Decode  $b$ 
  Replace completed sentences with new sentence from queue
end while

```

Figure 7: Decoding for top-up batching

## 4 Experimental Setup

We trained a sequence-to-sequence, encoder-decoder NMT system similar to that described in Sennrich et al. (2016). This uses recurrent neural networks with gated recurrent units. The input and output vocabulary size were both set to 85,000 subwords using byte-pair encoding (BPE) to adjust the vocabulary to the desired size. The hidden layer dimensions was set to 512.

For inference, we used and extend Amun (Junczys-Dowmunt et al., 2016), the fastest open-source inference engine we are aware of for the model used in this paper. We used a mini-batch of 128 sentences and maxi-batch of 1280 sentences, unless otherwise stated.

The hardware used in all experiments was an Nvidia GTX 1060 GPU on a host containing 8 Intel hypercores running at 2.8Ghz, 16GB RAM and SSD hard drive.

	Subset of training data	OpenSubtitles
# sentences	800k	200k
# subwords	5.8m	5.9m
Avg subwords/sent	7.3	29.7

Table 1: Corpora

Our training data consisted of the German-English parallel sentences from the Europarl corpus (Koehn, 2005). To test inference speed, we used two test sets with differing characteristics:

1. a subset of the training data, which contains mostly long sentences, and is, of course, in the same domain as the training data
2. a subset of the German-English data from the Open-Subtitles corpus, consisting of mostly short, out-of-domain sentences.

Table 1 gives further details of the test sets.

## 5 Results

## 6 Conclusion

### Acknowledgments

This work is sponsored by the Air Force Research Laboratory, prime contract FA8650-11-C-6160. The views and conclusions contained in this document are those of the authors and should not be interpreted as representative of the official policies, either expressed or implied, of the Air Force Research Laboratory or the U.S. Government.

### References

- Bahdanau, D., Cho, K., and Bengio, Y. (2014). Neural machine translation by jointly learning to align and translate. *CoRR*, abs/1409.0473.
- Cho, K., van Merriënboer, B., Gulcehre, C., Bahdanau, D., Bougares, F., Schwenk, H., and Bengio, Y. (2014). Learning phrase representations using rnn encoder–decoder for statistical machine translation. In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1724–1734. Association for Computational Linguistics.
- Devlin, J. (2017). Sharp models on dull hardware: Fast and accurate neural machine translation decoding on the CPU. In *Proceedings of the 2017 Conference on Empirical Methods in Natural Language Processing, EMNLP 2017, Copenhagen, Denmark, September 9-11, 2017*, pages 2820–2825.
- Guevara, M., Gregg, C., Hazelwood, K. M., and Skadron, K. (2009). Enabling task parallelism in the cuda scheduler.
- Jouppi, N. P., Young, C., Patil, N., Patterson, D., Agrawal, G., Bajwa, R., Bates, S., Bhatia, S., Boden, N., Borchers, A., Boyle, R., Cantin, P., Chao, C., Clark, C., Coriell, J., Daley, M., Dau, M., Dean, J., Gelb, B., Ghaemmaghami, T. V., Gottipati, R., Gulland, W., Hagmann, R., Ho, R. C., Hogberg, D., Hu, J., Hundt, R., Hurt, D., Ibarz, J., Jaffey, A., Jaworski, A., Kaplan, A., Khaitan, H., Koch, A., Kumar, N., Lacy, S., Laudon, J., Law, J., Le, D., Leary, C., Liu, Z.,

- Lucke, K., Lundin, A., MacKean, G., Maggiore, A., Mahony, M., Miller, K., Nagarajan, R., Narayanaswami, R., Ni, R., Nix, K., Norrie, T., Omernick, M., Penukonda, N., Phelps, A., Ross, J., Salek, A., Samadiani, E., Severn, C., Sizikov, G., Snelham, M., Souter, J., Steinberg, D., Swing, A., Tan, M., Thorson, G., Tian, B., Toma, H., Tuttle, E., Vasudevan, V., Walter, R., Wang, W., Wilcox, E., and Yoon, D. H. (2017). In-datacenter performance analysis of a tensor processing unit. *CoRR*, abs/1704.04760.
- Junczys-Dowmunt, M., Dwojak, T., and Hoang, H. (2016). Is neural machine translation ready for deployment? a case study on 30 translation directions. In *Proceedings of the 9th International Workshop on Spoken Language Translation (IWSLT)*, Seattle, WA.
- Kalchbrenner, N. and Blunsom, P. (2013). Recurrent continuous translation models. Seattle. Association for Computational Linguistics.
- Kim, Y. and Rush, A. M. (2016). Sequence-level knowledge distillation. In *Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing, EMNLP 2016, Austin, Texas, USA, November 1-4, 2016*, pages 1317–1327.
- Koehn, P. (2005). Europarl: A parallel corpus for statistical machine translation. In *Proceedings of the Tenth Machine Translation Summit (MT Summit X)*, Phuket, Thailand.
- Koehn, P. and Knowles, R. (2017). Six challenges for neural machine translation. In *Proceedings of the First Workshop on Neural Machine Translation*, pages 28–39, Vancouver. Association for Computational Linguistics.
- Lacey, G., Taylor, G. W., and Areibi, S. (2016). Deep learning on fpgas: Past, present, and future. *CoRR*, abs/1602.04283.
- Miciekevicius, P., Narang, S., Alben, J., Diamos, G. F., Elsen, E., Garcia, D., Ginsburg, B., Houston, M., Kuchaiev, O., Venkatesh, G., and Wu, H. (2017). Mixed precision training. *CoRR*, abs/1710.03740.
- Sennrich, R., Haddow, B., and Birch, A. (2016). Neural Machine Translation of Rare Words with Subword Units. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1715–1725, Berlin, Germany. Association for Computational Linguistics.